

JVM Performance and Tuning

Module Two Performance Metrics



Module Topics

- 1. Defining Performance**
- 2. Performance Metrics**
- 3. Performance Measurements**
- 4. Benchmarking and Profiling**

Some Caveats

- There are no “go faster” options for a JVM
- There are no list of tips and tricks to optimize a JVM
- The performance of code depends on many factors including:
 - The hardware the JVM is running on
 - The version of the JVM
 - How the code was designed and written
- Java design choices
 - Program safety over performance
 - Programmer productivity over performance

Java Design Decisions

- Designed for interactive applications
 - Generally the human is the rate limiting step
 - As long as it is “fast enough” then performance can be sacrificed for productivity
 - *Main cost is software development is programmer time, both writing code and debugging and maintaining code*
 - Interactive applications tend to require more updating and maintenance as requirements evolve and change
- Relies on a number of managed subsystems
 - e.g. Memory allocation and garbage collection
 - These subsystems introduce complexity into the running of the JVM
- Even measuring performance affects performance

Performance Metrics

- Depends on what is important to you
- Business level metrics
 - Ease of use
 - Customer loyalty
 - Improves business operations
 - “Apdex” index
 - *Measure of how long a transaction should take – industry standard*
- Application level metrics
 - Throughput measures – how many simultaneous transactions
 - Reliability, durability, loading etc
 - Quality levels, risk and error levels

Performance Metrics

- Depends on what is important to you.. cont
- Execution Metrics (what we are interested in)
 - Efficient use of hardware
 - CPU usage, memory usage, I/O bandwidth
- Source Code Metrics
 - Readability – ease of modification and changes
 - Effective and efficient design
 - Algorithm choice and implementation

Performance Metrics

- Application level metrics are often design or topology related
- Code metrics are often related to code design issues
- But both application and code issues affect execution metrics
- We are concerned with how to improve execution metrics if we assume there are no issues at the application topology and source code levels

Module Topics

1. Defining Performance
- 2. Performance Metrics**
3. Performance Measurements
4. Benchmarking and Profiling

Performance Metrics

- Standard metrics include:
 - Throughput
 - Latency
 - Capacity
 - Utilization
 - Efficiency
 - Scalability
 - Degradation

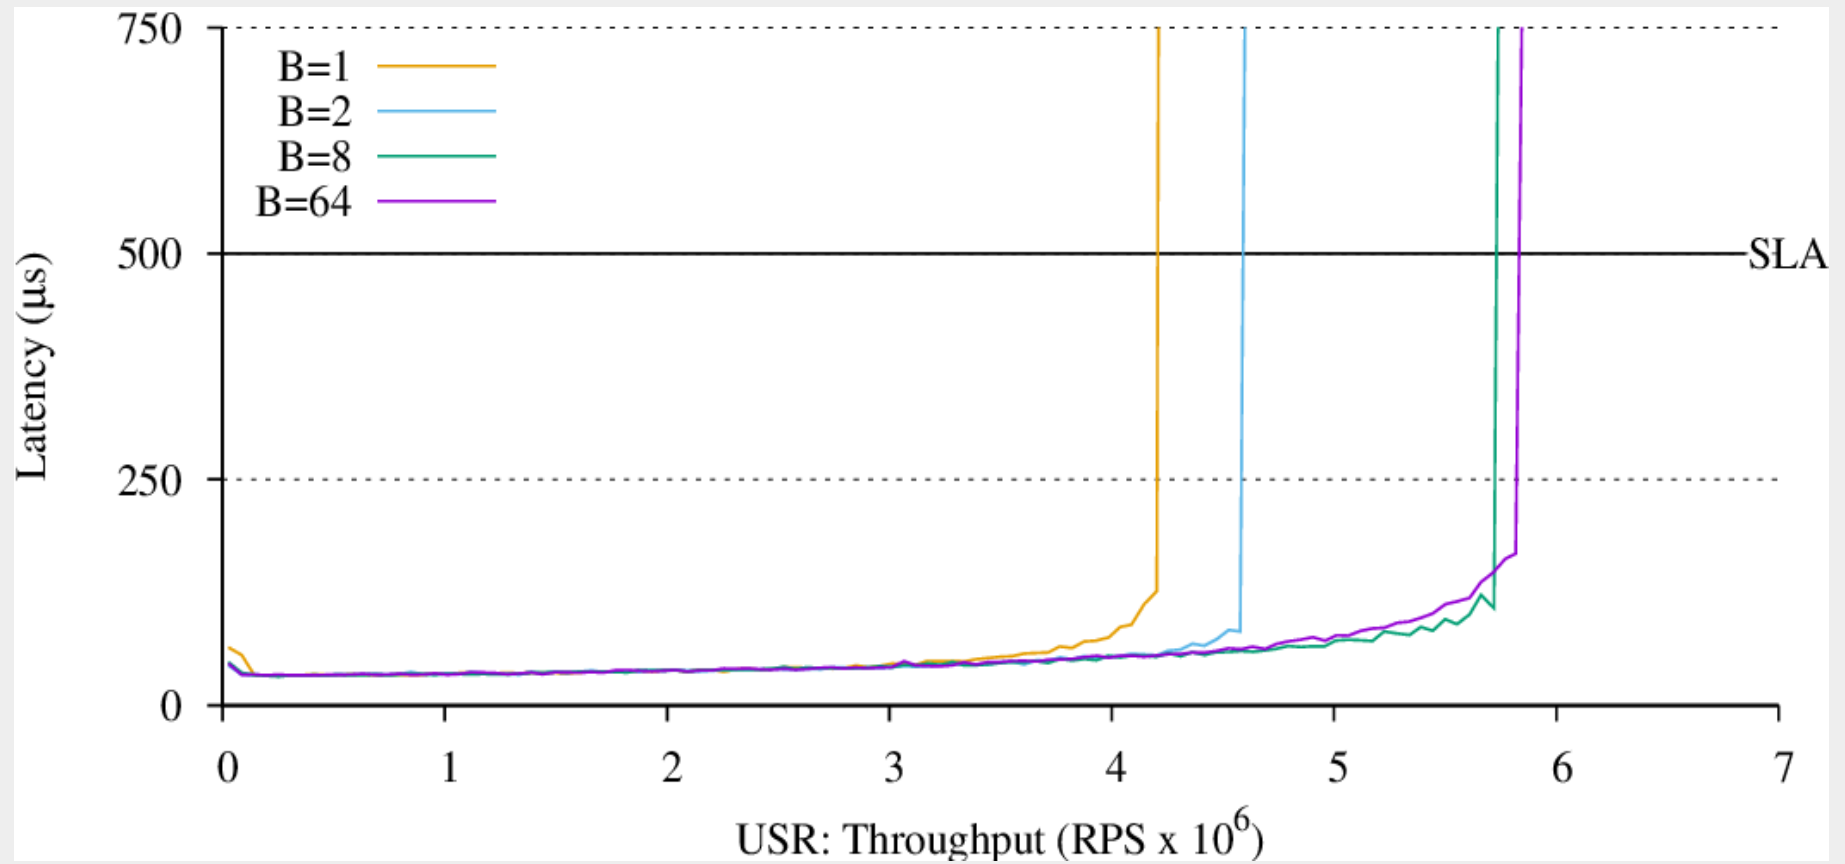
Throughput

- The rate of work a system can perform.
- Expressed as number of units of work in some time period
- Dependent on the platform where measured
- Dependent on how the unit of work is defined
- Often represented as number of complete transactions

Latency

- Time taken to process a single transaction
- Generally expressed as a start-to-end time.
- Dependent on workload
 - Expressed as a function of increasing workload
 - Increasing workload means less task access to resources
 - For example, browser latency increases the more tabs are open


Latency



Capacity

- Amount of work parallelism a system possesses
- The number of units of work that can be simultaneously ongoing in the system
- Capacity is related to throughput
 - The more capacity available implies more throughput possible
- Restaurants put more staff on at meal times
 - Increases restaurant capacity
 - Improves throughput = number of meals prepared
 - Reduces latency = the amount of time a customer has to wait


Capacity



Utilization

- Efficient use of a system's resources
- Computationally intensive tasks may have close to 100% CPU utilization but have little memory use
- A common problem today is:
 - CPU performance has increased dramatically
 - Memory performance has not
 - CPUs are idle waiting for memory to catch up moving data in and out of the CPU registers.

Utilization



Memory Access vs CPU Speed

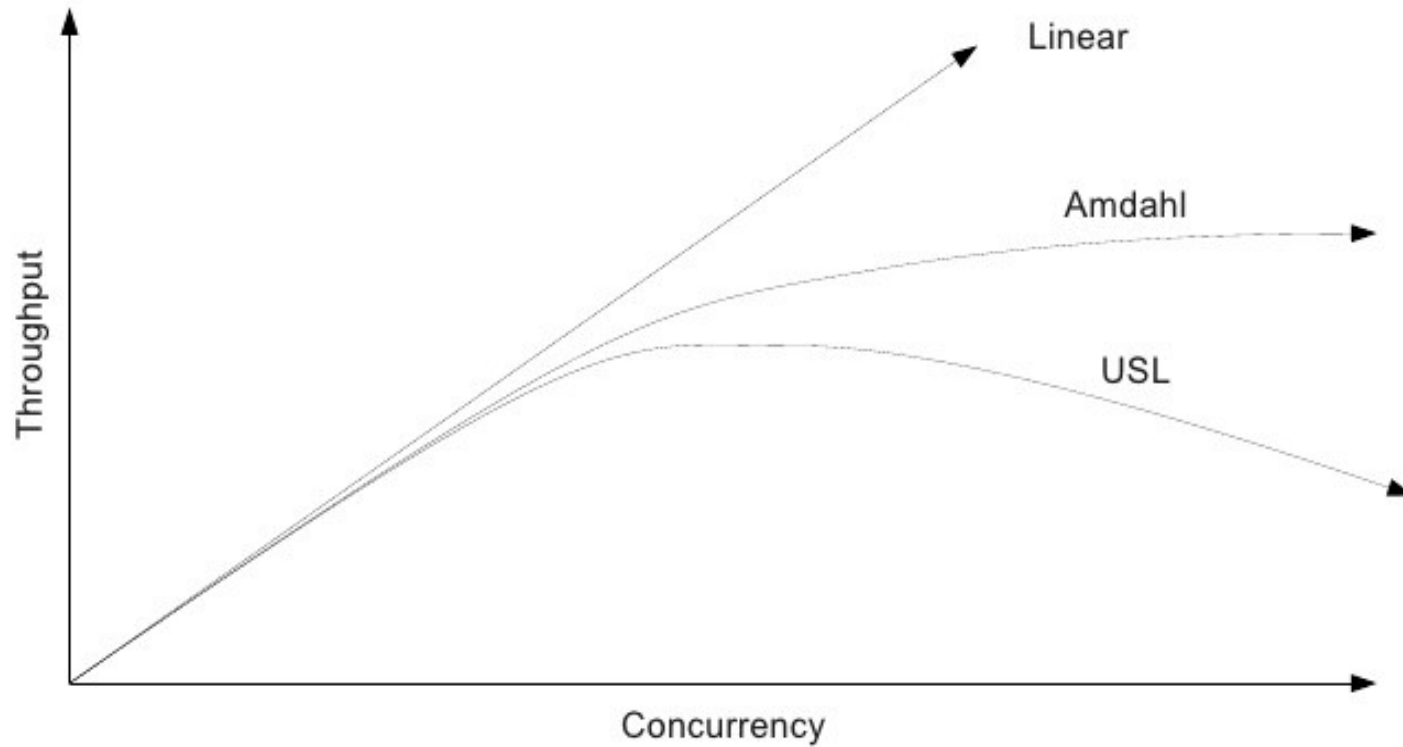
Efficiency

- Overall system efficiency is measured by
 - $(\text{throughput})/(\text{utilized resources})$
- Essentially making the most use of available resources
- Often used to compute costs of processing

Scalability

- Change in throughput as resources are added
- Ideally, throughput should increase with increase in capacity
- Scalability fails when increased capacity does not produce an increase in throughput

Universal Scalability Law



A top-down view of a wooden desk. In the top right, a portion of a silver laptop is visible, showing keys like 'tab', 'Q', 'W', 'E', 'caps lock', 'A', 'S', 'Z', 'fn', 'control', and 'option'. Below the laptop, a pair of black-rimmed glasses lies horizontally. To the right of the glasses is a white ceramic cup filled with dark coffee, with a yellow handle. In the top center, a small green succulent plant in a dark pot sits on the desk. The background is a light-colored wooden surface with a prominent grain pattern.

Module Topics

1. Defining Performance
2. Performance Metrics
- 3. Performance Measurements**
4. Benchmarking and Profiling

Performance Tests

- All tests, to be useful must be:
 - **Valid:** We have to actually test what we think we are testing
 - **Accurate:** The results are quantitatively accurate
 - **Reliable:** The results of the tests depend only on the tests and thing being tested
- These properties are true for all testing not just performance testing
 - Generally performance testing is statistical in nature
 - However it is difficult to identify all the contributing factors to the results of the tests
 - Good testing protocols need to be observed

Java Design

- Java performance philosophy
 - raw performance could be sacrificed if developer productivity benefited
 - provided performance was “good enough”
- Managed subsystems
 - low-level control exchanged for managed resources
 - eg. JVM memory management and garbage collection
 - Also provided a safer runtime “sandbox”

JVM Behaviour

- JVM runtime behaviour
 - Very complex – difficult to measure consistently
 - Observed measurements are often not normally distributed
 - Makes statistical analysis difficult – most statistical techniques assume underlying normal distributions
 - Critical measurements are easily missed via sampling
 - Outliers may be significant (“that one weird time that the app crashed”)
- Java performance measures can be misleading
 - We are trying to measure a highly complex set of subsystems
 - May even be non-linear in many aspects
 - *ie. small changes may produce large variations and vice versa*

The Experimenter Effect

- The actual measurement of performance affects performance
 - Measurement has overhead
 - Sampling measurements require resources
 - These resources are the same resources being measured
- This means that any sampling profile
 - Can only be regarded as approximate
 - And is not a true measure of performance.

A Basic Methodology

- Basic Steps
 - Define the desired outcome
 - Measure the existing system
 - Propose a modification that may produce the desired outcome
 - Apply the modification
 - Retest with the modification
 - Decide if the desired outcome has been achieved

Module Topics


1. Defining Performance
2. Performance Metrics
3. Performance Measurements
4. Benchmarking and Profiling

Benchmarking

- Compares two competing pieces of code
 - The code is executed a set number of times
 - Specific metrics are recorded each iteration
 - A final average results is computed
 - The process is repeated with the second piece of code
- Large numbers of repetitions are used
 - Evens out the variation we mentioned earlier
 - Not unusual to use hundreds of thousands of iterations

Profiling

- Looks for bottlenecks in a programs
 - Which methods are called
 - How long each method take to complete
 - Resource bottlenecks
 - *CPU bound – spends time using or waiting for CPU*
 - *IO bound – spends time waiting for or doing I/O*
 - *Memory bound – spends a lot of time paging memory*
 - Average values of data over multiple executions is often used

A top-down view of a wooden desk. In the top right, a portion of a silver laptop is visible, showing keys like 'ESC', '1', '2', '3', 'Q', 'W', 'E', 'tab', 'caps lock', 'A', 'S', 'Z', 'fn', 'control', and 'option'. Below the laptop, a pair of black-rimmed glasses lies horizontally. To the right of the glasses is a white ceramic cup filled with dark coffee, with a yellow handle. In the bottom right corner, the top edge of a black tablet is visible. At the top center, a small green succulent in a dark pot sits on the desk. A semi-transparent dark rectangle with the text 'Module End' is positioned on the left side of the desk.

Module End